

Draft — Abstract

Status: Voice pass 5 + IUI frame pass · 2026-07-14 (opens on LLM; light technical jargon; no 0%; no pilot)

Venue frame: ACM IUI 2027 — user control / steering of generative tutor replies under multi-turn drift

Abstract

LLM chat tutors are quickly becoming default study tools, specially for students with ADHD. Left unconstrained, these models drift in multi-turn dialogue: answers expand into verbose monologues, scannable structure collapses, and fluent generation can hallucinate detail the course materials never supplied. One-shot prompt instructions ("stay short," "use bullets") can help for a single turn. Across a conversation they fail to hold. For a learner who already struggles to re-orient after interruption, that drift is not cosmetic. It is an accessibility failure in the tutoring interface — and a control failure for anyone relying on prompt steering alone.

This paper treats ADHD-supportive response structure as a runtime enforcement problem for interactive LLM tutoring rather than a prompting trick. We run a controlled three-arm comparison on matched multi-turn tutoring probes: an unconstrained baseline tutor, the same tutor with an ADHD-supportive structure policy injected into the prompt only, and that policy plus a second-pass oversight checker that may revise the draft before it reaches the learner. Model, tools, and decoding settings stay fixed so the independent variable is response control alone. Without assistive control, the required scaffolds almost never hold. Prompt-only policy recovers most structural adherence. Second-pass oversight improves adherence further, but only modestly. The implication is sharp: accessibility that depends on perfect prompting will fail under ordinary dialogue. Reliable ADHD-supportive tutoring needs an enforceable structure layer — an intelligent control path over generative replies — that resists multi-turn drift.

Notes

- Opens on **LLM** as requested.
- IUI lean (light): "control failure," "interactive LLM tutoring," "intelligent control path over generative replies" — no invented topic laundry-list.
- Jargon used lightly: multi-turn, hallucinate, runtime enforcement, three-arm, independent variable, structural adherence, second-pass oversight, decoding settings.
- Still no EduAI, no Dean/Teacher/Student labels, no 0%, no pilot.

Voice-edit checklist

- ☐ Roommate can still follow the story
- ☐ Jargon supports the claim; doesn't replace it
- ☐ Hallucination = motivation, not measured DV